

## 1) Write operation

[root@swancloud ~]# rados bench -p pool1 50 write --no-cleanup

Maintaining 16 concurrent writes of 4194304 bytes for up to 50 seconds or 0 objects

Object prefix: benchmark\_data\_swancloud.novalocal\_98229

sec	Cur ops	started	finished	avg MB/s	cur MB/s	last lat	avg lat
0	0	0	0	0	-	0	
1	16	35	19	75.9789	76	0.659728	0.513125
2	16	50	34	67.9869	60	0.870065	0.6713
3	16	68	52	69.3224	72	1.18346	0.769565
4	16	77	61	60.9913	36	1.16245	0.818601
5	16	92	76	60.7918	60	1.56775	0.933608
6	16	98	82	54.6595	24	1.1883	0.958428
7	16	98	82	46.851	0	-	0.958428
8	16	98	82	40.9946	0	-	0.958428
9	16	106	90	39.9948	10.6667	3.99442	1.23474
10	16	125	109	43.5944	76	0.726977	1.40112
11	16	140	124	45.0851	60	0.977684	1.34673
12	16	152	136	45.3277	48	1.41589	1.33438
13	16	164	148	45.5328	48	0.988901	1.3383
14	16	173	157	44.8516	36	1.11751	1.33262
15	16	173	157	41.8614	0	-	1.33262
16	16	173	157	39.2452	0	-	1.33262
17	16	173	157	36.9366	0	-	1.33262
18	16	191	175	38.8841	18	0.626778	1.59996
19	16	206	190	39.9951	60	1.12779	1.55399

2015-10-14 19:31:40.745462min lat: 0.0234121 max lat: 4.73398 avg lat: 1.53213

sec	Cur ops	started	finished	avg MB/s	cur MB/s	last lat	avg lat
20	16	218	202	40.3951	48	1.13599	1.53213
21	16	233	217	41.3284	60	1.10702	1.50936
22	16	245	229	41.6314	48	1.22908	1.493
23	16	248	232	40.343	12	1.42215	1.49017
24	16	248	232	38.662	0	-	1.49017
25	16	251	235	37.5954	6	3.45147	1.51609
26	16	254	238	36.6109	12	4.21678	1.551
27	16	254	238	35.255	0	-	1.551
28	16	254	238	33.9959	0	-	1.551
29	16	269	253	34.8923	20	0.609106	1.78669
30	16	281	265	35.3291	48	1.17513	1.77776
31	16	293	277	35.7376	48	1.30729	1.75461
32	16	302	286	35.7457	36	1.2375	1.73844
33	16	311	295	35.7533	36	1.40236	1.73168
34	16	311	295	34.7017	0	-	1.73168
35	16	314	298	34.0531	6	3.22691	1.74781
36	16	329	313	34.7736	60	0.619422	1.80693
37	16	344	328	35.4552	60	1.2757	1.77764
38	16	356	340	35.7852	48	1.14409	1.75772
39	16	371	355	36.4059	60	1.07176	1.73486

2015-10-14 19:32:00.747810min lat: 0.0234121 max lat: 7.09035 avg lat: 1.71405

sec	Cur ops	started	finished	avg MB/s	cur MB/s	last lat	avg lat
40	16	383	367	36.6956	48	0.962997	1.71405
41	16	386	370	36.0932	12	1.22002	1.71032
42	16	386	370	35.2338	0	-	1.71032
43	16	389	373	34.6935	6	3.72288	1.72704
44	16	401	385	34.9958	48	3.6718	1.78922

45	16	413	397	35.2846	48	1.333	1.78643
46	16	422	406	35.3001	36	1.46305	1.78029
47	16	434	418	35.5702	48	1.18591	1.77123
48	16	446	430	35.829	48	1.27788	1.75932
49	16	452	436	35.5876	24	1.32821	1.75501
50	16	452	436	34.8758	0	-	1.75501
51	16	453	437	34.2704	2	3.47876	1.75895

Total time run: 51.294348

Total writes made: 453

Write size: 4194304

Bandwidth (MB/sec): 35.326

Stddev Bandwidth: 25.1211

Max bandwidth (MB/sec): 76

Min bandwidth (MB/sec): 0

Average Latency: 1.80626

Stddev Latency: 1.33042

Max latency: 7.09035

Min latency: 0.0234121

## 2) Sequential read operation

[root@swancloud ~]# rados bench -p pool1 50 seq

sec	Cur ops	started	finished	avg MB/s	cur MB/s	last lat	avg lat
0	0	0	0	0	-	0	
1	15	448	433	1731.77	1732	0.0382319	0.0362406

Total time run: 1.048970

Total reads made: 453

Read size: 4194304

Bandwidth (MB/sec): 1727.409

Average Latency: 0.0367782

Max latency: 0.0650652

Min latency: 0.00530613

## 3) Random read operation.

[root@swancloud ~]# rados bench -p pool1 50 rand

sec	Cur ops	started	finished	avg MB/s	cur MB/s	last lat	avg lat
0	1	1	0	0	-	0	
1	16	478	462	1847.59	1848	0.0348754	0.0339901
2	16	930	914	1827.45	1808	0.0341374	0.0347062
3	15	1381	1366	1820.89	1808	0.0364235	0.0349317
4	16	1862	1846	1845.59	1920	0.0266991	0.0345495
5	15	2460	2445	1955.6	2396	0.0280014	0.0326315
6	16	3032	3016	2010.28	2284	0.0280201	0.0317576
7	15	3582	3567	2037.92	2204	0.0287014	0.0313376
8	15	4011	3996	1997.66	1716	0.0372803	0.0319605
9	16	4434	4418	1963.23	1688	0.0378059	0.032527
10	15	4975	4960	1983.67	2168	0.0273686	0.0322181
11	15	5533	5518	2006.22	2232	0.0296165	0.0318572
12	15	6085	6070	2023.01	2208	0.0301363	0.0315942
13	15	6626	6611	2033.84	2164	0.0279156	0.0314325
14	15	7192	7177	2050.26	2264	0.0288929	0.0311815

15	15	7652	7637	2036.22	1840	0.0369566	0.0313912
16	16	8082	8066	2016.19	1716	0.0370977	0.0317035
17	15	8511	8496	1998.75	1720	0.0365232	0.0319845
18	16	8951	8935	1985.25	1756	0.0336392	0.0322059
19	15	9378	9363	1970.85	1712	0.0400492	0.0324389

2015-10-14 19:36:37.623782min lat: 0.00650516 max lat: 0.0521495 avg lat: 0.0326625

sec	Cur ops	started	finished	avg MB/s	cur MB/s	last lat	avg lat
20	15	9805	9790	1957.7	1708	0.033422	0.0326625
21	15	10269	10254	1952.83	1856	0.0339326	0.0327463
22	15	10730	10715	1947.87	1844	0.0348146	0.0328288
23	15	11205	11190	1945.78	1900	0.0341631	0.0328675
24	15	11664	11649	1941.19	1836	0.0356863	0.0329434
25	15	12118	12103	1936.15	1816	0.0340524	0.0330325
26	15	12551	12536	1928.29	1732	0.0364487	0.0331653
27	15	12981	12966	1920.57	1720	0.0379527	0.0332992
28	15	13407	13392	1912.83	1704	0.0374026	0.0334348
29	16	13830	13814	1905.07	1688	0.0340799	0.0335732
30	15	14294	14279	1903.56	1860	0.035742	0.0336004
31	15	14748	14733	1900.72	1816	0.0356741	0.0336513
32	16	15203	15187	1898.06	1816	0.0349786	0.033698
33	15	15648	15633	1894.6	1784	0.0374461	0.03376
34	16	16081	16065	1889.69	1728	0.0383952	0.0338471
35	15	16504	16489	1884.15	1696	0.0375721	0.0339484
36	15	16953	16938	1881.7	1796	0.0345443	0.0339948
37	16	17413	17397	1880.45	1836	0.0319224	0.0340178
38	16	17894	17878	1881.59	1924	0.0335153	0.033997
39	15	18362	18347	1881.44	1876	0.0340039	0.0340013

2015-10-14 19:36:57.627095min lat: 0.00650516 max lat: 0.0589415 avg lat: 0.0340448

sec	Cur ops	started	finished	avg MB/s	cur MB/s	last lat	avg lat
40	15	18808	18793	1879	1784	0.0345414	0.0340448
41	15	19261	19246	1877.36	1812	0.0374004	0.0340744
42	15	19713	19698	1875.7	1808	0.0356107	0.0341055
43	16	20167	20151	1874.21	1812	0.0353037	0.0341322
44	16	20627	20611	1873.43	1840	0.0344918	0.0341473
45	16	21084	21068	1872.41	1828	0.0346931	0.0341659
46	15	21545	21530	1871.87	1848	0.0343361	0.0341764
47	15	22007	21992	1871.36	1848	0.0336817	0.0341862
48	15	22468	22453	1870.79	1844	0.0342109	0.0341974
49	16	22930	22914	1870.24	1844	0.0344018	0.0342066
50	15	23367	23352	1867.87	1752	0.034168	0.0342468

Total time run: 50.030907  
Total reads made: 23367  
Read size: 4194304  
Bandwidth (MB/sec): 1868.205

Average Latency: 0.0342492  
Max latency: 0.0589415  
Min latency: 0.00650516

#### 4) Sequential read operation (10)

```
[root@swancloud ~]# rados bench -p pool1 10 seq
sec Cur ops  started finished avg MB/s  cur MB/s last lat  avg lat
0    1    1    0    0    0    -    0
Total time run:    0.970286
Total reads made:  425
Read size:        4194304
Bandwidth (MB/sec): 1752.060

Average Latency:    0.0363265
Max latency:        0.068168
Min latency:        0.00592239
```

#### 5) Random read operation (10)

```
[root@swancloud ~]# rados bench -p pool1 10 rand
sec Cur ops  started finished avg MB/s  cur MB/s last lat  avg lat
0    1    1    0    0    0    -    0
1   16   455   439 1755.61  1756 0.0388162 0.0356987
2   15   911   896 1791.65  1828 0.0388266 0.0353851
3   16  1395  1379 1838.3   1932 0.0407589 0.0345615
4   15  1802  1787 1786.68  1632 0.0382464 0.0356441
5   15  2214  2199 1758.89  1648 0.0399048 0.0362412
6   16  2634  2618 1745.04  1676 0.0390581 0.0365522
7   16  3109  3093 1767.14  1900 0.0328712 0.0361261
8   16  3595  3579 1789.21  1944 0.0336504 0.0356899
9   15  4082  4067 1807.26  1952 0.0332372 0.0353451
10  16  4577  4561 1824.11  1976 0.031427  0.03502
Total time run:    10.029614
Total reads made:  4577
Read size:        4194304
Bandwidth (MB/sec): 1825.394

Average Latency:    0.0350297
Max latency:        0.0575722
Min latency:        0.00637411
```

	50 Seq. read op	50 Random read	10 Seq. read op	10 Random read
B.W. (Mb/sec)	1727.409	1868.205	1752.060	1825.394
Avg. Latency	0.0367782	0.0342492	0.0363265	0.0350297
Total time run	1.048970	50.030907	0.970286	10.029614
Total reads made	453	23367	425	4577

## 1) Local disk benchmark

Always use the oflag=direct in order to use direct I/O. Because the system maintains a page cache to improve I/O performance. Every single write operations to the storage system are considered completed after the data has been copied to the page cache. The page cache is copied to permanent storage (hard drive disk) using the system call fsync(2).

```
[root@swancloud my-cluster]# dd if=/dev/zero of=here bs=1G count=1 oflag=direct
1+0 records in
1+0 records out
1073741824 bytes (1.1 GB) copied, 11.3945 s, 94.2 MB/s
[root@swancloud my-cluster]#
```

2)

```
[root@swancloud ~]# sudo bonnie++ -s 8192 -r 4096 -u root -d /mnt/ -m BenchClient
Using uid:0, gid:0.
Writing a byte at a time...done
Writing intelligently...done
Rewriting...done
Reading a byte at a time...done
Reading intelligently...done
start 'em...done...done...done...done...done...
Create files in sequential order...done.
Stat files in sequential order...done.
Delete files in sequential order...done.
Create files in random order...done.
Stat files in random order...done.
Delete files in random order...done.
Version 1.96 -----Sequential Output----- --Sequential Input- --Random-
Concurrency 1 -Per Chr- --Block-- -Rewrite- -Per Chr- --Block-- --Seeks--
Machine Size K/sec %CP K/sec %CP K/sec %CP K/sec %CP K/sec %CP /sec %CP
BenchClient 8G 1988 98 109317 6 111792 4 4244 99 7245719 99 3876 19
Latency 4072us 111us 179us 2063us 244us 2393us
Version 1.96 -----Sequential Create----- -----Random Create-----
BenchClient -Create-- --Read--- -Delete-- -Create-- --Read--- -Delete--
files /sec %CP /sec %CP /sec %CP /sec %CP /sec %CP /sec %CP
16 23410 26 +++++ + 31123 28 19735 22 +++++ + 27677 29
Latency 341us 61us 613us 4723us 29us 318ms
1.96,1.96,BenchClient,1,1445254220,8G,,1988,98,109317,6,111792,4,4244,99,7245719,99,3876,19,16,,,,,2
3410
,26,+++++,++++,31123,28,19735,22,+++++,++++,27677,29,4072us,111us,179us,2063us,244us,2393us,341us,
61us,613 us,4723us,29us,318ms
```